The Molecular Architecture of Discovery: From Atomic Fundamentals to Agentic AI in 2026 Pharmaceutical Prediction

The progression of human understanding from the fundamental constituents of matter to the autonomous orchestration of molecular design represents the most significant paradigm shift in scientific history. As of 2026, the pharmaceutical industry has transitioned from a traditional Research and Development (R&D) model to a predictive Research and Prediction (R&P) framework, where the "brute force" methods of the past have been supplanted by rational, AI-driven design.[1] This evolution is rooted in the realization that biological systems and the diseases that afflict them are, at their core, complex chemical signaling networks governed by the behavior of atoms and the electrons that define their interactions.[2]

The Foundation of Matter and the Cosmic Elemental Forge

Matter constitutes the physical substance of the universe, encompassing everything from the vast nebulae of deep space to the intricate cellular machinery of the human body.[2] Matter is characterized by its existence in distinct states, each defined by the kinetic energy and arrangement of its constituent particles. In the 2026 landscape of drug discovery, understanding these states is foundational to comprehending how drug molecules behave in different biological environments.

The Dynamics of Physical States

Matter traditionally exists in four primary states, and the transitions between them are governed by thermodynamic processes. In the context of pharmacology, these states influence drug formulation, solubility, and delivery mechanisms.

State of Matter	Transition Process (To)	Transition Process (From)	Description
Solid	Melting (to Liquid)	Freezing (from Liquid)	Particles are tightly packed in a fixed

			structure.
Liquid	Vaporization (to Gas)	Condensation (from Gas)	Particles are close but can move past each other.
Gas	Ionization (to Plasma)	Deionization (from Plasma)	Particles are far apart and move independently.
Plasma	—	—	Highly energized, ionized gas with free electrons.
Solid (direct)	Sublimation (to Gas)	Deposition (from Gas)	Phase change bypassing the liquid state.

2

These physical transitions are not merely abstract concepts; they are the basis for techniques like lyophilization (freeze-drying), which utilizes sublimation to stabilize volatile drug compounds.[2] Furthermore, the ionization of gases into plasma represents a high-energy state that, while rare in biological systems, provides insights into the high-energy environments where elements are forged.

Elements as the Pure Substances of the Universe

At the core of matter are elements—pure chemical substances that cannot be simplified into other substances by chemical means.[2] To date, scientists have identified 118 elements, 98 of which occur naturally on Earth.[2] The organization of these elements into the Periodic Table is not a mere exercise in memorization but a structural map of chemical behavior based on atomic architecture.[2]

The distribution of elements is starkly different between the cosmic scale and the terrestrial environment. In the vast majority of the universe, Hydrogen and Helium reign supreme,

accounting for 98% of the element space.[2] On Earth, the composition reflects a process of planetary differentiation where heavier elements concentrated in the crust and core.

Element	Universe Abundance (%)	Earth's Crust Abundance (%)	Primary Biological Role
Hydrogen (H)	75.0%	0.14%	Component of water and organic molecules.
Helium (He)	23.0%	Trace	No known biological role.
Oxygen (O)	Trace	47.0%	Essential for respiration; major crustal mass.
Silicon (Si)	Trace	28.0%	Major component of minerals; trace in body.
Aluminum (Al)	Trace	8.0%	Industrial metal; no biological role.
Iron (Fe)	Trace	5.0%	Oxygen transport in blood (hemoglobin).

2

While Oxygen and Silicon dominate the Earth's crust, forming the silicate minerals that comprise our rocks and soil, the human body is a specialized subset of this elemental pool, optimized for the chemical complexity of life.[5]

The Bio-Elemental Blueprint: Chemistry of Life

Life requires a highly specific set of elemental building blocks to maintain the complex signaling and structural integrity necessary for survival. Of the 118 known elements, only 25 are considered essential for human life.[2] These elements are categorized by their abundance and the criticality of their physiological functions.

The "Big Four" and Bulk Elements

Four elements alone account for approximately 96% of the mass of most living organisms.[2] These "bulk elements" form the fundamental structure of every protein, lipid, carbohydrate, and nucleic acid in the human body.

Element	Body Weight (%)	Symbol	Biological Importance
Oxygen	65.0%	O	Primary component of water and cellular respiration.
Carbon	18.5%	C	The backbone of all organic chemistry and life.
Hydrogen	9.5%	H	Essential for water and acid-base balance.
Nitrogen	3.3%	N	Found in all proteins and DNA/RNA bases.

2

Carbon's unique ability to form four stable covalent bonds allows for the creation of long, complex chains and rings that serve as the structural framework for biological molecules. Nitrogen is indispensable for the amino acids that build proteins and the nitrogenous bases that encode genetic information.[9]

Macrominerals and the Phosphorus Bottleneck

Beyond the primary four, six macrominerals are required in significant quantities (grams per

day) to maintain homeostasis, provide essential ions in body fluids, and form structural components.[11]

Macrominerals	**Body Weight (%)**	**Key Functions**
Calcium (Ca)	1.5%	Bone structure, muscle function, nerve signaling.
Phosphorus (P)	1.0%	Energy transfer (ATP), DNA backbone, bone health.
Potassium (K)	0.4%	Nerve function, osmotic balance, heart rhythm.
Sulfur (S)	0.3%	Found in amino acids like cysteine and methionine.
Sodium (Na)	0.2%	Fluid balance, transmission of nerve impulses.
Chlorine (Cl)	0.2%	Electrolyte balance, stomach acid (HCl) production.
Magnesium (Mg)	0.1%	Co-factor for 300+ enzymes, energy metabolism.

2

Phosphorus represents a unique "chemical bottleneck" in biology. While it accounts for 1.1% of the human body, it makes up only 0.105% of the Earth's crust.[5] This disproportionate need versus availability makes Phosphorus a limiting factor in many biological environments. It is a

crucial constituent of ATP (Adenosine Triphosphate), the universal energy currency of cells, and the phosphodiester bonds that link the genetic building blocks of RNA and DNA.[5]

The Critical Role of Trace Elements

Trace elements comprise less than 0.01% of body weight, yet their absence can lead to severe physiological malfunctions or death.[2] These elements often act as catalysts or structural centers for enzymes and vitamins.

- **Iron (Fe):** Essential for oxygen transport; its deficiency leads to anemia.[9]
- **Zinc (Zn):** Vital for immune function, wound healing, and cognitive processes; over a billion people are estimated to be at risk of Zinc deficiency.[12]
- **Iodine (I):** Required by the thyroid gland to produce hormones that regulate metabolism.[8]
- **Chromium (Cr):** In its Cr(III) form, it enhances insulin function, though its Cr(VI) form is highly toxic and carcinogenic.[12]
- **Lithium (Li):** Although not traditionally listed as essential, it is widely used to treat mood disorders and has been suggested to have longevity-enhancing properties.[10]

Maintaining a precise balance of these elements is critical, as excessive amounts can be toxic, while deficiencies disrupt the body's chemical equilibrium.[8]

Atomic Structure: The Smallest Unit of Chemical Identity

The atom is the smallest unit of an element that retains all its chemical properties.[2] To understand modern drug discovery, one must look beyond the macroscopic level into the subatomic architecture where drugs interact with their biological targets.

The Nuclear Core and the Electron Cloud

Atoms are predominantly empty space—over 99.94% of an atom's mass is concentrated in a microscopic nucleus at its center.[2] This nucleus consists of positively charged protons and neutral neutrons. Surrounding the nucleus is a diffuse "electron cloud" containing negatively charged electrons.[2]

The identity of an element is defined by its **Atomic Number**, which is the total count of protons in its nucleus.[2] Altering this number fundamentally changes the element. For example, adding a single proton to a Carbon atom (Atomic Number 6) transforms it into a Nitrogen atom (Atomic Number 7).[2]

The **Mass Number** represents the sum of protons and neutrons:

$$Mass\ Number = (\#protons) + (\#neutrons)$$

[2]

While the proton count is fixed for each element, the number of neutrons can vary, creating isotopes. Isotopes of the same element share chemical properties but differ in physical stability.

Carbon Isotope	Protons	Neutrons	Mass Number	Application
Carbon-12	6	6	12	Most common, stable form (98.9%).
Carbon-13	6	7	13	Stable, used in NMR spectroscopy.
Carbon-14	6	8	14	Unstable, used in radioactive dating.

[2]

The Electronic Basis of Pharmaceutical Action

A central tenet for students of drug discovery is that chemistry is "all about electrons".[2] The distribution and behavior of electrons in the outer shells of atoms determine the shape, charge, and reactivity of molecules. In the pharmaceutical context, a disease is often a malfunction in chemical signaling.[2] To "fix" this malfunction, a drug molecule—the "key"—must bind to a biological target—the "lock"—through a series of electronic interactions.[2]

Drug-target binding is determined by two primary factors:

1. **Shape:** The 3-D conformation of the molecule must fit into the "pocket" or "canyon" of the receptor or enzyme.[13]
2. **Charge:** Electrostatic attractions between the drug and its target pull the molecules

together and hold them in place.[13]

Interaction Type	Relative Strength	Character	Mechanism
Covalent Bond	Extremely Strong	Irreversible	Mutual sharing of an electron pair.
Ionic Interaction	Strong	Reversible	Attraction between opposite charges (cation/anion).
Hydrogen Bond	Moderate	Reversible	Attraction between H and electronegative atoms (N, O).
Van der Waals	Weak	Reversible	Forces from temporary shifting electron densities.

13

Most drug-target interactions rely on a combination of multiple weak, reversible forces. This allows the drug to bind, initiate a signaling change, and then depart.[14] However, certain toxic agents, like nerve gases, form irreversible covalent bonds with enzymes, leading to permanent inactivation and severe biological harm.[13]

The Pharmaceutical Imperative: Traditional Drug Discovery

Drug discovery is the systematic process of identifying chemical compounds that can safely and effectively treat disease. In the biological context, a target is usually a protein or gene associated with a specific ailment.[2] Traditional drug development, however, has historically been a slow, costly, and high-risk endeavor.

The Lifecycle and the "Wall" of Attrition

The traditional journey of a drug from discovery to market typically exceeds a decade and costs between $1 billion and $2.6 billion.[1] Despite these resources, roughly 90% of drugs that enter clinical trials fail to achieve approval.[1]

Stage of Discovery	Key Objective	Primary Reason for Failure
Hit Discovery	Identify molecules that interact with the target.	Low selectivity or weak binding.
Lead Optimization	Refine hit molecules for potency and safety.	Emerging toxicity or poor metabolic profile.
Preclinical Testing	Validate in cell lines and animal models.	Discrepancy between animal and human results.
Phase I Trials	Assess safety in humans.	Unanticipated toxicity.
Phase II/III Trials	Assess efficacy and side effects.	Lack of clinical efficacy (40-50% of failures).

[1]

The primary causes of failure include a lack of clinical efficacy—where the drug simply doesn't work in humans as it did in the lab—and high clinical toxicity.[1] Toxicity often arises from "off-target" effects, where a drug binds to unintended receptors, or "on-target" inhibition in vital organs where the target protein performs essential functions.[1]

ADMET: The Five Pillars of Drug Viability

To succeed, a molecule must possess more than just target affinity; it must have a favorable ADMET profile.[17]

- **Absorption:** Can the drug enter the bloodstream (e.g., across the intestinal wall)? [18]
- **Distribution:** Can it reach the target tissue, including crossing the blood-brain barrier

(BBB) if necessary? [18]
- **Metabolism:** How is the drug broken down by the body, particularly by liver enzymes? [17]
- **Excretion:** How quickly is the drug removed from the system? [17]
- **Toxicity:** Does the drug cause harm to cells, organs, or the overall system? [17]

Traditional drug discovery often identifies ADMET issues too late in the process. This has created a massive opportunity for Deep Learning to accelerate the screening of these properties in the "in silico" (digital) phase before lab testing begins.[2]

The Deep Learning Revolution in Molecular Design

The 2026 landscape of drug discovery is defined by the integration of Deep Learning (DL) models that can predict molecular behavior with high accuracy. These models accelerate the transition from a "hit" to a "lead" compound, reducing the time and capital wasted on candidates destined to fail.[1]

Variational Autoencoders (VAE) and Chemical Generation

One of the most transformative tools in modern chemistry is the Variational Autoencoder (VAE). VAEs are used to "generate" new candidate molecules that have never existed before.[2] Unlike standard models that simply copy data, a VAE learns the underlying patterns and "grammar" of chemical structures, represented as SMILES strings.[2]

The VAE architecture consists of an **Encoder** and a **Decoder**:

1. The Encoder compresses a discrete molecular structure into a continuous "latent space" representation.[20]
2. In this latent space, molecules are represented as coordinates. By moving through this space, researchers can find "neighboring" molecules that are slight variations of the original lead compound.[2]
3. The Decoder then translates these coordinates back into a valid chemical structure.[2]

This allows for the exploration of a vast, untapped chemical space, enabling researchers to design molecules with desired biological activities and optimized pharmacological properties.[20]

Drug-Target Interaction (DTI) and DeepPurpose

Predicting how well a drug "key" fits into a protein "lock" is the role of Drug-Target Interaction (DTI) models.[2] The DeepPurpose library is a comprehensive deep learning framework that has standardized this process.[2]

DeepPurpose takes two primary inputs: the drug's SMILES string and the protein's amino acid sequence.[22] It utilizes diverse encoders to map these inputs into vector representations.

Encoder Type	Used For	Strengths
MLP (on Fingerprints)	Drugs	Uses traditional descriptors like Morgan or RDKit.
CNN	Drugs/Proteins	Captures local patterns in strings and sequences.
RNN	Drugs/Proteins	Captures sequential dependencies in data.
Transformer	Drugs/Proteins	Captures long-range relationships using attention mechanisms.

[22]

By concatenating these learned embeddings and passing them through a decoder, DeepPurpose predicts the **binding affinity** (e.g., IC_{50}).[22] This allows researchers to rank millions of virtual molecules and select only the most promising candidates for synthesis.[26]

Deep2Lead: A Unified Lead Optimization Interface

Deep2Lead is an example of a distributed application that bridges the gap between complex DL models and bench scientists. It integrates a VAE for molecule generation and DeepPurpose for DTI prediction into a single interface.[2]

In a landmark case study, Deep2Lead was used to optimize Imatinib (Gleevec), an anti-cancer drug targeting the BCR-ABL fusion protein.[2] By starting with Imatinib's structure and its target sequence, the system successfully generated a novel inhibitor that was predicted to be **10 times more potent** than the original drug.[2] This demonstrated the capability of AI to perform sophisticated lead optimization tasks without requiring the user to have extensive programming expertise.[2]

The 2026 Frontier: Agentic AI and Predictive Pipelines

The cutting edge of drug discovery in 2026 has moved beyond simple predictive models toward **Agentic AI**. These are autonomous systems capable of reasoning, planning, and executing multi-step research goals with minimal human oversight.[2]

The Shift from R&D to R&P (Research & Prediction)

The 2026 landscape is defined by the transition to Research and Prediction (R&P), where the predictive engine effectively replaces the screening lab.[1]

Feature	Traditional R&D	2026 Agentic R&P
Discovery Method	Brute force mass-screening.	Rational digital design/blueprinting.
AI Status	"Black Box" results.	"Glass Box" with transparent reasoning.
Cycle Time	Years for candidate nomination.	Months (or weeks) for nomination.
Focus	Manual data cross-checking.	High-level strategic review.

1

Agentic AI acts as a "digital detective," reasoning through noisy data to isolate the faint "biological signals" that cause rare diseases.[1] By modeling how proteins misbehave, these agents can nominate drug candidates with unprecedented speed, ensuring that only the highest-quality therapies enter the clinical value chain.[1]

Autonomous Screening and Synthesis

Agentic systems like "Coscientist" and "ChemCrow" have demonstrated the ability to autonomously plan chemical syntheses and orchestrate robotic laboratory equipment.[30]

- **In Silico Design:** Agents scan millions of molecules, predicting efficacy, toxicity, and

ADMET profiles simultaneously.[32]

- **Synthesis Planning:** Agents determine the optimal timing for critical steps in chemical synthesis, such as solvent switches, replacing countless manual iterations.[33]
- **Closed-Loop Discovery:** In a "self-driving laboratory," the AI generates a hypothesis, plans the experiment, executes it via robotics, analyzes the results, and refines the hypothesis without human intervention between steps.[28]

Real-World Deployment in Pharma Giants

Major pharmaceutical companies have already embraced this shift. Johnson & Johnson (J&J) utilizes agentic systems for autonomous drug discovery and the optimization of chemical synthesis steps.[33] Pfizer has deployed agentic AI to predict patient enrollment rates in clinical trials and identify potential dropout risks, significantly reducing trial timelines and costs.[33] These systems have moved from experimental pilots to being deeply embedded in daily operations at an industrial scale.[29]

Model Context Protocol (MCP) and Data Interoperability

One of the most critical technical enablers of the 2026 landscape is the **Model Context Protocol (MCP)**. This open standard allows AI models to connect seamlessly to massive, real-time biological and chemical databases.[2]

Solving the "N x M" Integration Problem

Historically, connecting an AI application to multiple databases required a custom connector for every single source—a non-scalable "N x M" problem.[35] MCP standardizes this connection, acting like a "USB-C port for AI systems".[35] It allows a single client (like an AI agent) to interact with multiple servers (databases like ChEMBL or UniProt) through a unified grammar.[34]

MCP Implementation	Purpose/Focus	Benefit
GEOmcp	Gene Expression Omnibus	Autonomous discovery of genomic datasets.
STRING MCP	Protein-Protein Interactions	Automated network construction and analysis.

ChEMBL-MCP	Chemical Informatics	Tools for molecular property prediction.
BioContextAI	Community Registry	Findable, metadata-rich MCP catalogues.

36

By using MCP, AI agents can dynamically retrieve and process up-to-date scientific data during inference.[36] This ensures that the AI is grounded in scientifically validated information, reducing "hallucinations" and making the research process more reproducible.[38] For example, an agent can autonomously identify a mouse brain dataset in GEO, extract relevant genes, and immediately transfer them to STRING for network analysis, performing a multi-day bioinformatics task in minutes.[37]

The Future of Compliance and Regulatory Pipelines

The impact of agentic AI extends beyond the lab and into the regulatory "war room." By 2026, compliance has become autonomous and proactive.[1]

AI-Ready Submissions and Accelerated Queries

The manual marathon of preparing global submissions is being replaced by multi-agent systems that connect all data sources and manage documentation in real-time.[1]

- **Synchronous Documentation:** The submission package is effectively built while the drug is being developed, eliminating inconsistencies that cause regulatory delays.[1]
- **Traceable Reasoning:** Unlike older "Black Box" models, agentic systems provide a transparent audit trail of their reasoning, linking every fact in a submission back to its verifiable source.[1]
- **Accelerated Query Response:** Addressing complex regulatory queries from the FDA or EMA, which once took weeks of manual effort, can now be accomplished in a two-day turnaround using fully cited, agent-generated responses.[1]

Regulatory Guardrails for 2026

The shift toward autonomous agents in regulated environments has been supported by the publication of joint principles for good AI practice by the FDA and EMA on January 14, 2026.[29] This transatlantic alignment on AI validation gives pharmaceutical companies the confidence to move AI into GxP-related territories, such as automated pharmacovigilance and dossier

preparation, knowing they meet rigorous audit and traceability standards.[29]

Synthesis and Conclusion: The R&P Paradigm

The transition from fundamental chemistry to the 2026 agentic landscape represents a total reimagining of how humanity creates medicine. Chemistry, the science of matter and electrons, provides the "grammar" for this new era. The Periodic Table and atomic theory are no longer just academic concepts but the parameters within which deep learning models and autonomous agents operate to solve the malfunctions of disease.[2]

By 2026, the industrialization of "in silico" trials and the use of AI-generated "digital twins" have drastically reduced Phase II/III trial durations and recruitment hurdles for rare diseases.[29] It is expected that AI-native discovery platforms will be responsible for over 60% of new drugs entering Phase I trials by the end of 2026.[29]

The "R&P" shift has fundamentally altered the role of the scientist. By automating the administrative burden, the "brute force" screening, and the clerical cross-checking of data, agentic AI has freed researchers to focus on high-level strategic review and high-impact scientific breakthroughs.[1] The goal of this modern tutorial is to ensure that students understand this bridge: that a mastery of "how atoms bind" is the prerequisite for commanding the AI agents that will discover the cures of tomorrow.[2] The molecular architecture of discovery is now a unified pipeline—from the subatomic electron cloud to the global autonomous research network—dedicated to the preservation of human health through the predictive power of science.

Works cited

1. Agentic AI: The shift from R&D to R&P will deliver the first predictive ..., accessed May 8, 2026, https://pharmaphorum.com/digital/agentic-ai-shift-rd-rp-will-deliver-first-predictive-drug-pipeline-2026
2. Drug Discovery and DL 101.pdf
3. Acta Interactive Archive Vol 51 - Sigma-Aldrich, accessed May 8, 2026, https://www.sigmaaldrich.com/deepweb/assets/sigmaaldrich/product/documents/580/438/acta-51-1-ms.pdf
4. CHEM20212 Computational Chemistry II: Fundamentals of electronic structure theory - Manchester Molecular Magnetism Group, accessed May 8, 2026, http://www.molmag.manchester.ac.uk/software/CHEM20212_2020_1.pdf
5. 2.1 The Elements | The Basics of General, Organic, and Biological Chemistry, accessed May 8, 2026, https://courses.lumenlearning.com/suny-monroecc-orgbiochemistry/chapter/the-elements/
6. Visualizing the abundance of elements in the Earth's crust - The World Economic Forum, accessed May 8, 2026, https://www.weforum.org/stories/2021/12/abundance-elements-earth-crust/

7. What Is the Most Abundant Element on Earth? Not Carbon - Science | HowStuffWorks, accessed May 8, 2026, https://science.howstuffworks.com/most-abundant-element-on-earth.htm
8. The essential elements - Science Learning Hub, accessed May 8, 2026, https://www.sciencelearn.org.nz/resources/1728-the-essential-elements
9. What are the 25 elements in the human body? - askIITians, accessed May 8, 2026, https://www.askiitians.com/forums/11-grade-biology-others/what-are-the-25-elements-in-the-human-body-25_476645.htm
10. Biological roles of the elements - Wikipedia, accessed May 8, 2026, https://en.wikipedia.org/wiki/Biological_roles_of_the_elements
11. Essential Elements for Life, accessed May 8, 2026, https://saylordotorg.github.io/text_general-chemistry-principles-patterns-and-applications-v1.0/s05-08-essential-elements-for-life.html
12. Our essential elements | Feature from King's College London, accessed May 8, 2026, https://www.kcl.ac.uk/our-essential-elements
13. Content Background: How Does a Drug Interact With Its Target? It's ..., accessed May 8, 2026, https://sites.duke.edu/thepepproject/module-5-why-do-plants-make-drugs-for-humans/content-background-how-does-a-drug-interact-with-its-target-its-all-in-the-chemistry/
14. Receptor-drug interaction - Philadelphia University, accessed May 8, 2026, https://www.philadelphia.edu.jo/academics/balakumar/uploads/4_Receptor-drug%20interactions.pdf
15. Drug–Receptor Interactions - Basicmedical Key, accessed May 8, 2026, https://basicmedicalkey.com/drug-receptor-interactions-2/
16. Machine Learning in Drug Discovery: A Review - PMC, accessed May 8, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC8356896/
17. ADMET Predictor® - Simulations Plus - Machine Learning- ADMET property prediction, accessed May 8, 2026, https://www.simulations-plus.com/software/admetpredictor/
18. Calculate ADMET properties - YouTube, accessed May 8, 2026, https://www.youtube.com/watch?v=TqSBXSkSxqA
19. Deep Batch Active Learning for Drug Discovery - eLife, accessed May 8, 2026, https://elifesciences.org/reviewed-preprints/89679v1
20. Molecule Generation of Drugs Using VAE - Atlantis Press, accessed May 8, 2026, https://www.atlantis-press.com/article/126002037.pdf
21. Evolution of Molecular Generation Models in Drug Discovery - ACS Publications, accessed May 8, 2026, https://pubs.acs.org/doi/10.1021/acs.jmedchem.5c03310
22. DeepPurpose: a deep learning library for drug–target interaction ..., accessed May 8, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC8016467/
23. DeepPurpose: A Deep Learning Library for Drug-Target Interaction Prediction arXiv:2004.08919v3 [cs.LG] 9 Dec 2020, accessed May 8, 2026, https://arxiv.org/pdf/2004.08919
24. Drug-protein interaction prediction via variational autoencoders and attention mechanisms, accessed May 8, 2026,

https://pmc.ncbi.nlm.nih.gov/articles/PMC9614151/
25. (PDF) DeepPurpose: A Deep Learning Library for Drug-Target Interaction Prediction, accessed May 8, 2026, https://www.researchgate.net/publication/347571227_DeepPurpose_A_Deep_Learning_Library_for_Drug-Target_Interaction_Prediction
26. DeepPurpose: a Deep Learning Based Drug Repurposing Toolkit - ICML Workshop on Computational Biology, accessed May 8, 2026, https://icml-compbio.github.io/icml-website-2020/2020/papers/WCBICML2020_paper_15.pdf
27. Agentic AI in healthcare: Types, trends, and 2026 forecast - Kellton, accessed May 8, 2026, https://www.kellton.com/kellton-tech-blog/agentic-ai-healthcare-trends-2026
28. FAQ: What You Need to Know About Agentic AI in Drug Discovery | BioPharm International, accessed May 8, 2026, https://www.biopharminternational.com/view/faq-what-you-need-to-know-about-agentic-ai-in-drug-discovery
29. 3 AI trends reshaping healthcare and life sciences in 2026 - Dataiku, accessed May 8, 2026, https://www.dataiku.com/stories/blog/healthcare-life-sciences-ai-trends-2026
30. Beyond SMILES: Evaluating Agentic Systems for Drug Discovery - arXiv, accessed May 8, 2026, https://arxiv.org/html/2602.10163v1
31. From Co-Pilot to Lab-Pilot: How Agentic AI is Redefining Chemical R&D - Cypris, accessed May 8, 2026, https://www.cypris.ai/insights/from-co-pilot-to-lab-pilot-how-agentic-ai-is-redefining-chemical-r-d
32. Artificial Intelligence (AI) in Drug Discovery: The Complete Guide (2026) - Bio-in-Tech, accessed May 8, 2026, https://bio-in-tech.com/blog/ai-in-drug-discovery-complete-guide-2026/
33. How Agentic AI is transforming life sciences discovery and operations - Medable Inc., accessed May 8, 2026, https://www.medable.com/knowledge-center/how-agentic-ai-is-transforming-life-sciences-discovery-and-operations
34. Model Context Protocol (MCP) Explained - Oracle, accessed May 8, 2026, https://www.oracle.com/database/model-context-protocol-mcp/
35. What is the Model Context Protocol (MCP)? - Databricks, accessed May 8, 2026, https://www.databricks.com/blog/what-is-model-context-protocol
36. How to Implement MCP Server in Drug Discovery AI - newline, accessed May 8, 2026, https://www.newline.co/@Dipen/how-to-implement-mcp-server-in-drug-discovery-ai--27c85713
37. MCPmed: a call for Model Context Protocol-enabled bioinformatics ..., accessed May 8, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12927880/
38. Connecting AI to biology: Model Context Protocol | EMBL, accessed May 8, 2026, https://www.embl.org/news/people-perspectives/connecting-ai-to-biology-model-context-protocol/